



2026-07-18-regression-baseline - v-pre-s1-studio
Sequence to read this file : Screenshot -> Details, Screenshot -> Details, ..




SS-1
This is what i see immediately after running the server




SS-2

After clicking on “New Idea”




SS-3

This is after running the idea




SS-4

This is the 1st artifact “Idea Intake” - Stage 0




SS-5

Click on “Actions” section



SS-6



Click on “history” section



SS-7

This is the Studio, reading state, An artifact is open and being viewed.

Below mentioned is the empty state, artifacts, etc



SS-7

Chrome

File Edit View History Bookmarks Profiles Tab Window Help

localhost

localhost:3000/improve

Open in app

Relaunch to update

Desk

Studio

Office

Settings

Stage 1 / 14

A voice-first AI assistant that seamlessly joins y

Running...

Nemotron 3 Super (Free)

Generating

Stage 1/14

Discovery

2 artifacts so far

"A voice-first AI assistant that ... this can take a minute or two on free models"

Stop

REFINE

Artifact Improvement Window

Improved 0 state

ARTIFACT GRAPH

EVENT FEED

REASONING CHAIN

15 nodes

0 state

0 improved

ARTIFACTS

2

SUGGESTED START

discovery

Start here - strongest foundation for the build package

Idea Intake

discovery

IMPROVEMENT INTENT

Describe what to improve - e.g. "Strengthen competitive moat with a defensible advantage and long-term implications"

Nemotron 3 Super (Free)

More concise

More strategic

More technical

More MVP-focused

Competitive moat

Market positioning

Long-term implications

Sharpen summary

Add evidence

Stronger recommendation

Improve structure

REFERENCE DOCS

Upload Reference Document

EXTENT

Light

Medium

Strong

SCOPE

Block

Stage

Project

Full artifact for this stage

Runtime Workflow

Click a stage node in the graph above

OR select from the artifact list on the left

Choose your model (8 providers via OpenRouter)

Describe your improvement intent or use a preset

Click IMPROVE NOW - this makes a real LLM call

Accept - artifact saved + state state propagates

Generate build package to send to any builder platform

Running Stage 1

Nemotron 3 Super (Free)

Stage 1 / 14

AI Meeting Assistant Rewrites

Chats are reviewed and used to improve Google AI Gemini's AI and can make mistakes, including about people. You can manage your activity, including info about your location. Learn more about your choices

rewrite this better:

Here are a few ways to rewrite your product description, depending on where you plan to use it:

Option 1: Clean & Professional (Best for Pitch Decks or Landing Pages)

A voice-first AI assistant that seamlessly joins your meetings—with or without you present. It captures full recordings and transcripts, generates real-time notes, tracks action items, and contextually answers questions on the spot using data from both the current meeting and your past conversations.

Option 2: Short & Punchy (Best for Taglines or Social Media)

An autonomous, voice-first meeting assistant that records, takes notes, and tracks tasks even if you can't make the call. Ask it anything mid-meeting, and it will instantly pull context from your

Sharing 'IdeaGate - AI-Native PM Operating Sys...'

Type @ to ask about a tab

Flash

Chrome

File Edit View History Bookmarks Profiles Tab Window Help

localhost

localhost:3000/improve

Open in app

Relaunch to update

Desk

Studio

Office

Settings

Stage 2 / 14

A voice-first AI assistant that seamlessly joins your meetings-with or without you present. It captures fu

Running...

Nemotron 3 Super (Free)

Generating

Stage 2/14

Problem Definition

3 artifacts so far

"A voice-first AI assistant that seamlessly joins your meetings-with or without you pr... this can take a minute or two on free models"

Stop

REFINE

Artifact Improvement Window

Improved 0 state

ARTIFACT GRAPH

EVENT FEED

REASONING CHAIN

15 nodes

0 state

0 improved

ARTIFACTS

3

Idea Intake

discovery

problem definition

IMPROVEMENT INTENT

Describe what to improve - e.g. "Strengthen competitive moat with a defensible advantage and long-term implications"

Nemotron 3 Super (Free)

More concise

More strategic

More technical

More MVP-focused

Competitive moat

Market positioning

Long-term implications

Sharpen summary

Add evidence

Stronger recommendation

Improve structure

REFERENCE DOCS

Upload Reference Document

EXTENT

Light

Medium

Strong

SCOPE

Block

Stage

Project

Full artifact for this stage

Current

2-problem-definition.md

Problem Definition: Voice-First AI Meeting Assistant

JTBD Analysis

Job Statement:

When professionals attend meetings (in-person, remote, or hybrid), they need to capture decisions, action items, and context accurately and instantly so they can stay accountable, follow up without delay, and retrieve information later without manual note-taking.

Desired Outcomes (based on discovery interviews):

Eliminate manual note-taking and post-meeting cleanup (target: reduce note-related work from 30-45 min/meeting to <5 min).

Ensure no action item is missed (target: >85 % recall of action items).

Enable instant retrieval of past meeting context via natural-language questions (target: answer relevance score >4/5).

Provide seamless integration with existing calendar and video-conferencing tools (zero-click join).

Guarantee privacy and control over meeting data (encryption, deletion, consent).

Customer Segments & Contexts (from Personas):

Busy Executives: strategy & update meetings, need rapid decision capture.

Project Managers: daily stand-ups, sprint planning, need action-item tracking.

Remote Team Members: asynchronous catch-up, need searchable record.

Administrative Assistants: board & all-hands meetings, need accurate minutes.

Current Alternatives & Their Shortcomings:

Manual note-taking: high cognitive load, error-prone.

Basic transcription tools (Otter, Fireflies): lack action-item detection, no contextual Q&A.

Sales-focused analytics (Gong): not suited for internal meetings.

Native AI notes (Teams Copilot): platform-locked, limited customization.

Running Stage 2

Nemotron 3 Super (Free)

Stage 2 / 14

AI Meeting Assistant Rewrites

Chats are reviewed and used to improve Google AI Gemini's AI and can make mistakes, including about people. You can manage your activity, including info about your location. Learn more about your choices

rewrite this better:

Here are a few ways to rewrite your product description, depending on where you plan to use it:

Option 1: Clean & Professional (Best for Pitch Decks or Landing Pages)

A voice-first AI assistant that seamlessly joins your meetings—with or without you present. It captures full recordings and transcripts, generates real-time notes, tracks action items, and contextually answers questions on the spot using data from both the current meeting and your past conversations.

Option 2: Short & Punchy (Best for Taglines or Social Media)

An autonomous, voice-first meeting assistant that records, takes notes, and tracks tasks even if you can't make the call. Ask it anything mid-meeting, and it will instantly pull context from your

Sharing 'IdeaGate - AI-Native PM Operating Sys...'

Type @ to ask about a tab

Flash

Chrome | File | Edit | View | History | Bookmarks | Profiles | Tab | Window | Help | Sat 3:16 PM

IdeaGate - AI-Native PM Op | localhost | Ask Gemini

localhost:3000/improve | Open in app | Relaunch to update

Desk

Studio

Office

Settings

Stage 2 / 14

A voice-first AI assistant that seamlessly joins your meetings-with or without you present. It captures fu

Running

Nemotron 3 Super (Free)

Stop

REFINE

ARTIFACT IMPROVEMENT TRACKING

Improved 0 state

MD

ARTIFACT GRAPH

EVENT FEED

REASONING CHAIN

15 nodes 0 state 0 improved

ARTIFACTS: 3

Idea Intake

discovery

problem definition

DOWNSTREAM: 12

solution design

mvp hypothesis

prd

ux design

validation

5. Trust & Privacy: Users distrust tools that cannot guarantee accuracy, security, and data control.

Assumption Map

Assumption Type	Description	Validation Method	Current Evidence	Risk if False
Desirability	Users will adopt an AI that joins meetings automatically and provides real-time notes + action-item sync.	Survey willingness-to-pay, prototype usability test.	14/20 willing to pay \$15-\$30/mo; 16/20 miss action items weekly.	Low adoption if users prefer manual notes or distrust AI.
Feasibility	Streaming ASR + intent extraction can achieve <2 s latency and >85 % action-item recall on diverse accents.	Technical spike measuring WDA and recall on internal dataset.	Early benchmarks show Whisper-large WDA ~4 % on clean audio; action-item extraction prototype ~78 % recall.	High latency or poor recall undermines core value prop.
Viability	Enterprise customers will pay for SOC 2/ISO 27001-compliant solution with per-meeting opt-in.	Pricing interviews, analyst TAM sizing.	14/20 interviewees indicated willingness to pay; security cited as deal-breaker.	Pricing too low or compliance gaps block enterprise sales.
Usability	Zero-click calendar join and native task-system sync will not disrupt existing workflows.	Usability testing with Outlook/Google Calendar and Jira/Asana adapters.	Positive feedback on seamless integration in discovery (theme: Seamless Integration).	Added friction leads to workarounds usage, reducing adoption.
Information	Contextual Q&A across past meetings creates a sticky knowledge-base effect.	Measure Q&A relevance and retention in pilot.	12/20 expressed desire to search past meetings; competitive gap identified.	If Q&A is inaccurate, retention benefit disappears.
Privacy	Users will accept cloud-based ASR if data is encrypted at rest/in-transit and deletable on request.	Privacy-concern survey, legal review.	Theme: Privacy & Control highlighted; 15/20 demand granular consent.	Lack of trust blocks adoption in regulated industries.

Key Risks to Mitigate:

- Action-item recall below target → improve entity-recognition models, incorporate domain-specific phrasing.
- Latency >2 s → optimize streaming pipeline, consider edge-buffering.
- Enterprise sales cycle length → prioritize SOC 2, ISO 27001, and offer on-prem/VPC deployment options.
- Privacy perception → transparent data-handling UI, independent audit reports.

Next Steps (Problem Definition Phase)

IMPROVEMENT INTENT

Describe what to improve - e.g., "Strengthen competitive moat with a defensible advantage and long-term implications"

Nemotron 3 Super (Free)

Generate

From: Fine-tune context generation

PRESETS

More concise

More strategic

More technical

More MVP-focused

Competitive moat

Market positioning

Long-term implications

Sharpen summary

Add evidence

Stronger recommendation

Improve structure

REFERENCE DOCS

Upload Reference Document

PDF DOC XLS MD Multiple allowed

EXTENT

LIGHT

MEDIUM

STRONG

Rewrite sections

SCOPE

BLOCK

STAGE

PROJECT

Full artifact for this stage

Generate

Nemotron 3 Super (Free)

Generate

From: Fine-tune context generation

BUILD DESTINATION

Running Stage 2

Nemotron 3 Super (Free)

Stage 2 / 14

Chrome | File | Edit | View | History | Bookmarks | Profiles | Tab | Window | Help | Sat 3:17 PM

IdeaGate - AI-Native PM Op | localhost | Ask Gemini

localhost:3000/improve | Open in app | Relaunch to update

Desk

Studio

Office

Settings

Stage 4 / 14

A voice-first AI assistant that seamlessly joins your meetings-with or without you present. It captures fu

Running

Nemotron 3 Super (Free)

Stop

REFINE

ARTIFACT IMPROVEMENT TRACKING

Improved 0 state

MD

ARTIFACT GRAPH

EVENT FEED

REASONING CHAIN

3 nodes 0 state 0 improved

ARTIFACTS: 3

Idea Intake

discovery

problem definition

DOWNSTREAM: 12

solution design

mvp hypothesis

prd

ux design

validation

Problem Definition: Voice-First AI Meeting Assistant

JTBD Analysis

Job Statement:

When professionals attend meetings (in-person, remote, or hybrid), they need to capture decisions, action items, and context accurately and instantly so they can stay accountable, follow up without delay, and retrieve information later without manual note-taking.

Desired Outcomes (based on discovery interviews):

- Eliminate manual note-taking and post-meeting cleanup (target: reduce note-related work from 30-45 min/meeting to <5 min).
- Ensure no action item is missed (target: >85 % recall of action items).
- Enable instant retrieval of past meeting context via natural-language questions (target: answer relevance score ≥4/5).
- Provide seamless integration with existing calendar and video-conferencing tools (zero-click join).
- Guarantee privacy and control over meeting data (encryption, deletion, consent).

Customer Segments & Contexts (from Personas):

- Busy Executives: strategy & update meetings, need rapid decision capture.
- Project Managers: daily stand-ups, sprint planning, need action-item tracking.
- Remote Team Members: asynchronous catch-up, need searchable record.
- Administrative Assistants: board & all-hands meetings, need accurate minutes.

Current Alternatives & Their Shortcomings:

- Manual note-taking: high cognitive load, error-prone.
- Basic transcription tools (Otter, Fireflies): lack action-item detection, no contextual Q&A.
- Sales-focused analytics (Gong): not suited for internal meetings.
- Native AI notes (Teams Copilot): platform-locked, limited customization.

Problem Framing

Problem Statement (One-Pager):

Knowledge workers lose an average of 30-45 minutes per meeting to note-taking, action-item tracking, and post-meeting

IMPROVEMENT INTENT

Describe what to improve - e.g., "Strengthen competitive moat with a defensible advantage and long-term implications"

Nemotron 3 Super (Free)

Generate

From: Fine-tune context generation

PRESETS

More concise

More strategic

More technical

More MVP-focused

Competitive moat

Market positioning

Long-term implications

Sharpen summary

Add evidence

Stronger recommendation

Improve structure

REFERENCE DOCS

Upload Reference Document

PDF DOC XLS MD Multiple allowed

EXTENT

LIGHT

MEDIUM

STRONG

Rewrite sections

SCOPE

BLOCK

STAGE

PROJECT

Full artifact for this stage

Generate

Nemotron 3 Super (Free)

Generate



From: Fine-tune context generation



BUILD DESTINATION



Running Stage 4


Nemotron 3 Super (Free)

Stage 4 / 14












Immediately after running a new idea -
Pages needs contextual re-work (with regards to what user should see)



After running the new idea still above to see the old artifacts



The new artifact 1st doc appeared in esk but the Studio doesn't have auto refresh where the old artifact goes and the new artifacts autonomously populate as and when generated



Then click on the studio and now can see the 2 docs on the left

The screenshot shows the IdeaGate AI interface in a web browser. The left sidebar has a 'Studio' tab selected. The main content area displays a document titled 'Stage: Idea Intake - Autonomous Voice-First Meeting Assistant'. The document includes an introduction, an NSM Lens section, key features, and benefits. On the right, there are sections for 'Improvement Intent', 'Presets', 'Reference Docs', 'Extent', 'Scope', and 'Build Destination'. The interface is dark-themed and includes a top navigation bar with various icons and a bottom status bar.

This is where i stop the test

The screenshot shows the OpenRouter logs page. The left sidebar has a 'Logs' tab selected. The main content area displays a table of logs. The table has columns for Date, Model, Provider, App ID, Input, Output, Cost, Usage Type, Speed, Finish Reason, and API Key. The logs show various API requests and responses, including those from the IdeaGate application.


Date	Model	Provider	App ID	Input	Output	Cost	Usage Type	Speed	Finish Reason	API Key
Jul 18, 03:28 PM	Llama 3.3 70B Instruct	DeepInfra	IdeaGate PMOS	1,743 tok	1,413 tok	\$ 0.000626	Credits	17.6 tok/s	stop	IdeaGate-Only
Jul 18, 03:23 PM	Llama 3.3 70B Instruct	DeepInfra	IdeaGate PMOS	2,222 tok	1,118 tok	\$ 0.00058	Credits	12.8 tok/s	stop	IdeaGate-Only
Jul 18, 03:22 PM	Llama 3.3 70B Instruct	DeepInfra	IdeaGate PMOS	1,802 tok	1,145 tok	\$ 0.000547	Credits	12.6 tok/s	stop	IdeaGate-Only
Jul 18, 03:22 PM	Llama 3.3 70B Instruct	Parasail	IdeaGate PMOS	1,805 tok	1,108 tok	\$ 0.000951	Credits	138.7 tok/s	stop	IdeaGate-Only
Jul 18, 03:21 PM	Llama 3.3 70B Instruct	DeepInfra	IdeaGate PMOS	877 tok	799 tok	\$ 0.000343	Credits	27.1 tok/s	stop	IdeaGate-Only
Jul 18, 03:21 PM	Llama 3.3 70B Instruct	Nebius Token Factory	IdeaGate PMOS	401 tok	485 tok	\$ 0.000246	Credits	38.4 tok/s	stop	IdeaGate-Only
Jul 18, 03:18 PM	Nemotron 3 Super (free)	NVIDIA	IdeaGate PMOS	1,642 tok	3,421 tok	\$ 0.00	Credits	79.2 tok/s	stop	IdeaGate-Only
Jul 18, 03:17 PM	Nemotron 3 Super (free)	NVIDIA	IdeaGate PMOS	2,144 tok	4,469 tok	\$ 0.00	Credits	102.5 tok/s	stop	IdeaGate-Only
Jul 18, 03:16 PM	Nemotron 3 Super (free)	NVIDIA	IdeaGate PMOS	616 tok	1,879 tok	\$ 0.00	Credits	140.8 tok/s	stop	IdeaGate-Only
Jul 18, 03:16 PM	Nemotron 3 Super (free)	NVIDIA	IdeaGate PMOS	2,907 tok	2,557 tok	\$ 0.00	Credits	119.9 tok/s	stop	IdeaGate-Only
Jul 18, 03:16 PM	Nemotron 3 Super (free)	NVIDIA	IdeaGate PMOS	2,145 tok	2,508 tok	\$ 0.00	Credits	109.8 tok/s	stop	IdeaGate-Only
Jul 18, 03:15 PM	Nemotron 3 Super (free)	NVIDIA	IdeaGate PMOS	2,140 tok	3,253 tok	\$ 0.00	Credits	119.7 tok/s	stop	IdeaGate-Only
Jul 18, 03:15 PM	Nemotron 3 Super (free)	NVIDIA	IdeaGate PMOS	2,622 tok	2,268 tok	\$ 0.00	Credits	62.0 tok/s	stop	IdeaGate-Only
Jul 18, 03:14 PM	Nemotron 3 Super (free)	NVIDIA	IdeaGate PMOS	2,140 tok	2,431 tok	\$ 0.00	Credits	73.9 tok/s	stop	IdeaGate-Only
Jul 18, 03:13 PM	Nemotron 3 Super (free)	NVIDIA	IdeaGate PMOS	2,762 tok	2,395 tok	\$ 0.00	Credits	63.1 tok/s	stop	IdeaGate-Only

This is what the open router looks like
API activity logs

The screenshot shows a web browser window at `localhost:3000/improve`. The page displays a document titled "Stage: Idea Intake - Autonomous Voice-First Meeting Assistant". The document content includes an introduction, a section on NSM Lenses, and key features. The interface also shows a console with network messages and a settings panel.

Developer mode

The screenshot shows the same web browser window, but with the console open, displaying a React DevTools error message: "Warning: Error: <circle> attribute r: Expected length, 'undefined'." The main content area displays the same document as the previous screenshot.



Mission control/analytics